



Discrete Event Simulation

Modelação e Desempenho de Redes e Serviços

Prof. Amaro de Sousa (asou@ua.pt)

DETI-UA, 2024/2025

Discrete event simulation

A discrete event simulation models the operation of a system whose state changes with events that happen in discrete time instants:

- each event forces a change of either a system state value and/or of a variable value;
- between consecutive events, the system remains in the same state;
- thus, the simulation can directly jump in time from one event to the next event.

Elements of a discrete event simulator:


- (1) State variables: variables that describe the state of the system at any time instant
- (2) Statistical counters: variables that store the statistical data required to determine the performance of the system
- (3) Simulation clock: variable indicating the current simulated time instant
(simulated time \neq computation time)
- (4) Events: types of occurrences that change either the system state and/or the statistical counters
- (5) Event list: list of future events, their time instants and associated parameters

Besides these elements, additional supporting variables might be required.

Basic structure of a discrete event simulator

A discrete event simulator is mainly composed by the following steps:

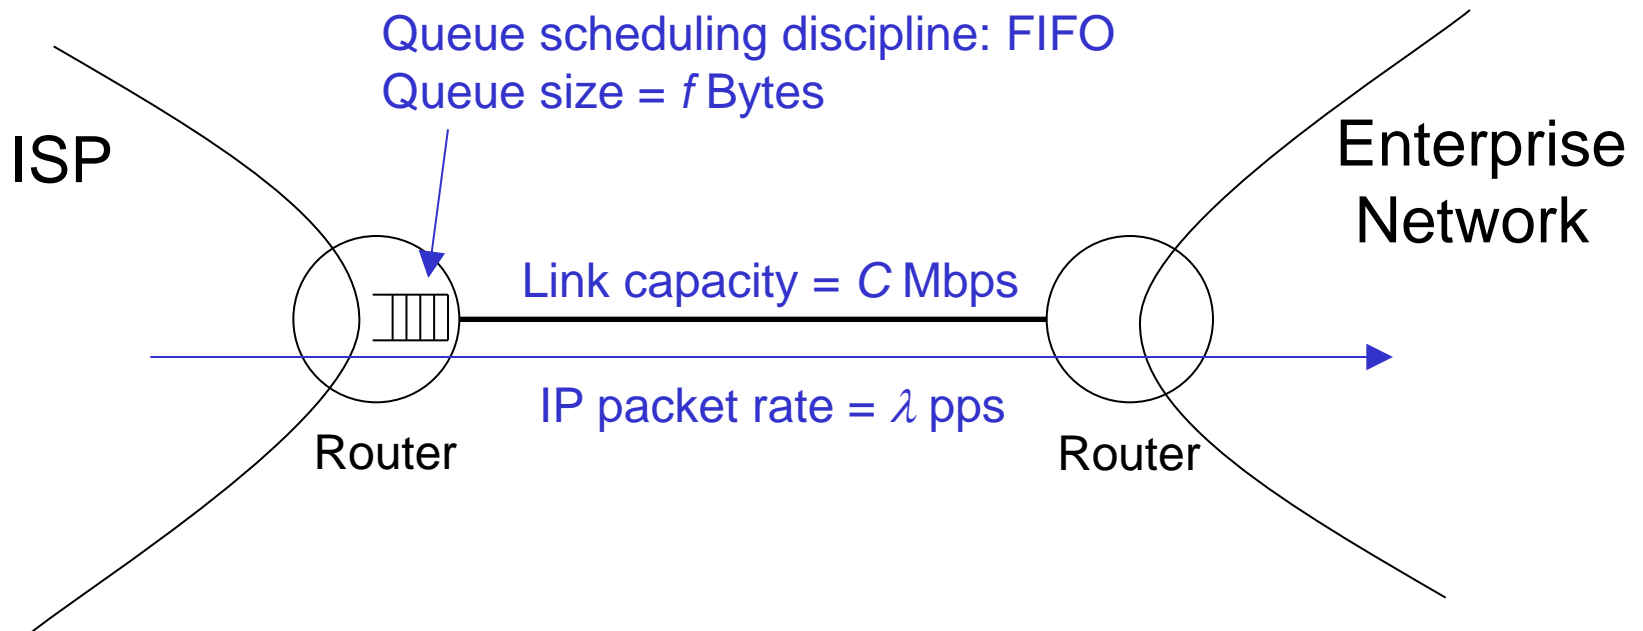
- (1) Initialization of the state variables, the statistical counters and the event list with the first event(s).
- (2) Determination of the next event from the event list.
- (3) Update of the simulation clock to the time instant of the next event and removal of the event from the event list.
- (4) Execution of all actions associated to the event (generation of new events, update of state variables and/or statistical counters).
- (5) Check if the simulation must end; if not, return to Step (2).
- (6) Update of the statistical counters (if needed) and determination of the performance parameters.



Example: performance of a downstream link from an ISP to a client enterprise

Consider the event driven simulation of a point-to-point IP link between an enterprise router and its Internet Service Provider (ISP).

Let us consider the downstream direction, *i.e.*, from ISP to the company (usually, the direction with highest traffic load).



Example: performance of a downstream link from an ISP to a client enterprise

Input parameters of simulation:

- λ – packet rate, in packets per second (pps)
- C – connection capacity, in Mbps
- f – queue size, in Bytes
- P – total number of transmitted packets of a simulation run

Stopping criteria of simulation:

Time instant when the link finishes the transmission of the P^{th} packet

Performance parameters to be estimated by the simulation:

- PL – Packet Loss (%)
- APD – Average Packet Delay (milliseconds)
- MPD – Maximum Packet Delay (milliseconds)
- TT – Transmitted Throughput (Mbps)

Example: performance of a downstream link from an ISP to a client enterprise

Events:

- ARRIVAL – the arrival of a packet
DEPARTURE – the end of transmission of a packet

State variables:

STATE – binary variable indicating if the connection is free (i.e., not being used) or busy with the transmission of a packet

QUEUEOCCUPATION – occupation of the queue, in number of bytes, with the queued packets

QUEUE – matrix with (i) a number of rows equal to the number of queued packets and (ii) 2 columns where each column has the size and the arriving time instant of each packet in the queue

Example: performance of a downstream link from an ISP to a client enterprise

Statistical Counters:

TOTALPACKETS – number of packets arrived to the system

LOSTPACKETS – number of packets dropped due to buffer overflow

TRANSPACKETS – number of transmitted packets

TRANSBYTES – sum of the bytes of the transmitted packets

DELAYS – sum of the delays of the transmitted packets

MAXDELAY – maximum delay among all transmitted packets

Performance parameters (at the end of the simulation):

$$PL = 100 \times \text{LOSTPACKETS} / \text{TOTALPACKETS}$$

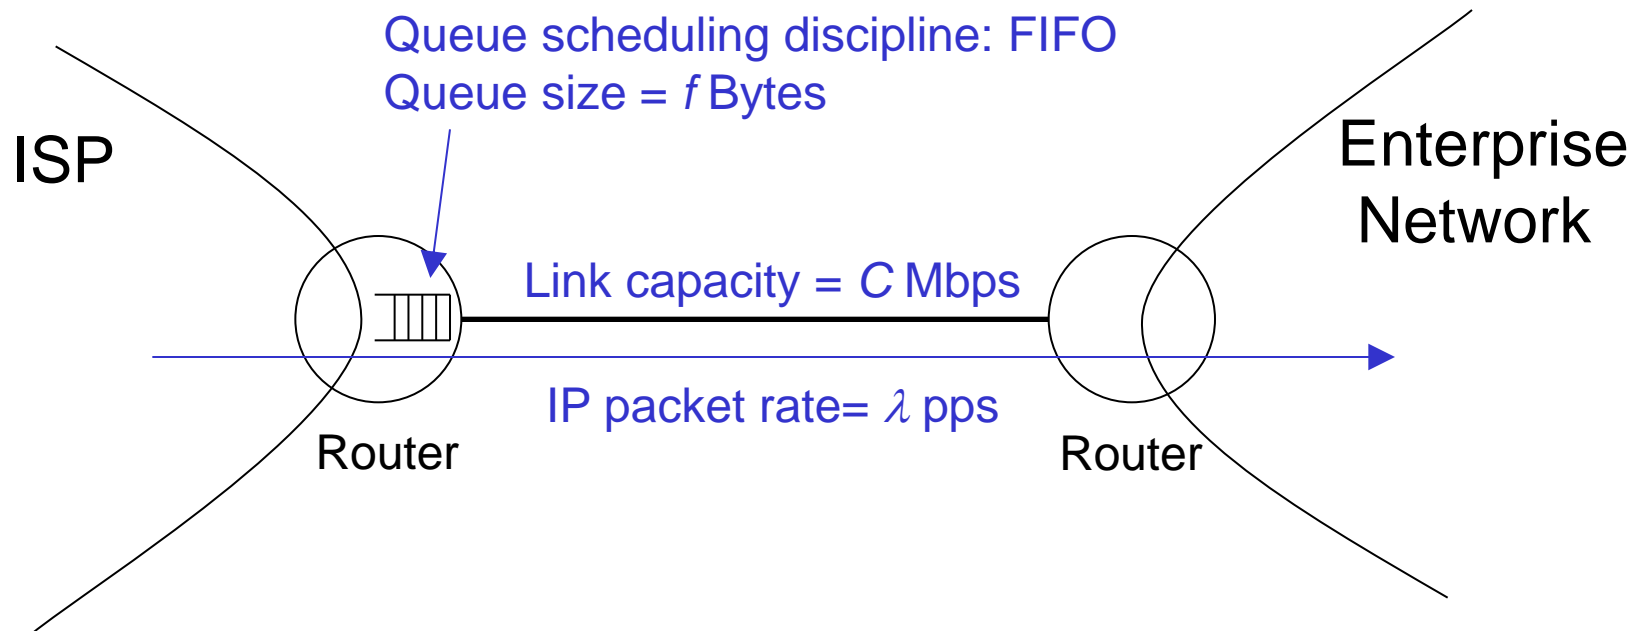
$$APD = 1000 \times \text{DELAYS} / \text{TRANSPACKETS}$$

$$MPD = 1000 \times \text{MAXDELAY}$$

$$TT = 10^{-6} \times \text{TRANSBYTES} \times 8 / \text{total simulated time}$$

Example: performance of a downstream link from an ISP to a client enterprise

Illustration of a simulation run in MATLAB



Generation of random numbers with a uniform distribution between 0 and 1

A Linear Congruential Generator (LCG) is an algorithm that yields a sequence of randomized numbers calculated with a linear equation.

The method represents one of the oldest and best-known pseudorandom number generator algorithms.

Generation method:

(1) Generate integer values Z_1, Z_2, \dots with the following recursive expression:

$$Z_i = (aZ_{i-1} + c) \pmod{m}$$

where m, a, c and Z_0 are non-negative integer parameters;

(2) Compute $U_i = Z_i / m$.

Seed of the generator



The values U_i seem to be real values uniformly distributed on the interval $[0,1]$

Example

Example: $Z_i = (5Z_{i-1} + 3)(\text{mod } 16)$

$$Z_0 = 7$$

| i | Z_i | U_i | i | Z_i | U_i |
|-----|-------|-------|-----|-------|-------|
| 0 | 7 | ---- | 10 | 9 | 0.563 |
| 1 | 6 | 0.375 | 11 | 0 | 0.000 |
| 2 | 1 | 0.063 | 12 | 3 | 0.188 |
| 3 | 8 | 0.500 | 13 | 2 | 0.125 |
| 4 | 11 | 0.688 | 14 | 13 | 0.813 |
| 5 | 10 | 0.625 | 15 | 4 | 0.250 |
| 6 | 5 | 0.313 | 16 | 7 | 0.438 |
| 7 | 12 | 0.750 | 17 | 6 | 0.375 |
| 8 | 15 | 0.938 | 18 | 1 | 0.063 |
| 9 | 14 | 0.875 | 19 | 8 | 0.500 |

- In this example, $m = 16$ and the algorithm repeats the generated numbers after 16 iterations (we say the generator has a period of 16).
- The random generator (function *rand*) of MATLAB has a period of $2^{31}-1$.

Generation of random numbers with other distributions

Discrete variables:

Consider the generation of a random variable that can be X_1, X_2, \dots, X_n .

Consider the probability of value X_i as $P(X = X_i) = f_i$.

Method:

- Split the interval $[0,1]$ in n intervals proportional to f_i , $i = 1 \dots n$
- Generate a uniformly distributed random value U in $[0,1]$
- Return X_i if U falls into the i -th interval

For example, the Bernoulli variable X with $p(0) = 1/4$ and $p(1) = 3/4$ can be randomly generated as:

(1) Generate $U \sim U(0,1)$

(2) If $U \leq 1/4$, return $X = 0$; otherwise, return $X = 1$

Generation of random numbers with other distributions

Discrete variables (MATLAB example)


```
x= [1 3 4 8 10];  
f= [0.1 0.2 0.2 0.4 0.1];
```

```
figure(1)  
stem(x,f)  
axis([-1 11 0 0.5])  
grid on
```

```
f_cum= [0 cumsum(f)]
```


```
a= zeros(1,100000);  
for it= 1:100000  
    a(it)= x(sum(rand()>f_cum));  
end
```

```
figure(2)  
hist(a,1:10)  
grid on
```



f_cum =

0.0 0.1 0.3 0.5 0.9 1.0



Generation of random numbers with other distributions

Generate a random packet size between 64 and 1518 bytes with the probabilities: 19% for 64 bytes, 23% for 110 bytes, 17% for 1518 bytes and an equal probability for all other values (i.e., from 65 to 109 and from 111 to 1517).

Custom MATLAB function:

```
function out= GeneratePacketSize()  
    aux= rand();  
    aux2= [65:109 111:1517];  
    if aux <= 0.19  
        out= 64;  
    elseif aux <= 0.19 + 0.23  
        out= 110;  
    elseif aux <= 0.19 + 0.23 + 0.17  
        out= 1518;  
    else  
        out = aux2(randi(length(aux2)));  
    end  
end
```

Generation of random numbers with other distributions

Continuous variables:

The most popular methods are based on the inverse of the cumulative distribution function (cdf).

Consider $F(X)$ as the cdf of a continuous random variable and $F^{-1}(U)$ as its inverse function.

Method:

(1) Generate $U \sim U(0,1)$

(2) Return $X = F^{-1}(U)$

For example, an exponential distributed random variable with average $1/\lambda$:


$$F(x) = \begin{cases} 1 - e^{-\lambda x}, & x \geq 0 \\ 0, & x < 0 \end{cases} \quad F^{-1}(U) = -\frac{1}{\lambda} \ln(U)$$

In MATLAB, use function *exprnd()*.

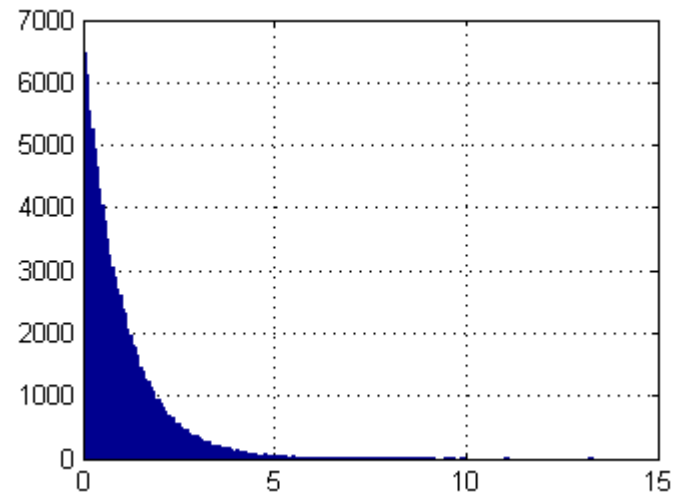
Generation of random numbers with other distributions

Exponential variable (MATLAB example)

```
lambda= 1  
x= 0:0.1:14;  
f=exppdf(x,1/lambda)  
figure(1)  
plot(x,f)  
grid on
```



```
a=exprnd(1/lambda,1,100000);  
figure(2)  
hist(a,200)  
grid on
```



Analysis of the results of a simulation

Consider X_1, X_2, \dots, X_n as the observations of independent and identically distributed (IID) random variables with average μ and finite variance σ^2 (for example, the results of different simulations of a given system).

The sample mean defined by
$$\bar{X}(n) = \frac{\sum_{i=1}^n X_i}{n}$$
 is an estimator for average μ .

The sample variance defined by
$$S^2(n) = \frac{\sum_{i=1}^n (X_i - \bar{X}(n))^2}{n-1}$$
 is an estimator for variance σ^2 .

The analysis of the results of a simulation is, usually, based on the Central Limit Theorem.

Analysis of the results of a simulation

Consider Z_n as a random variable given by:
$$Z_n = \frac{\bar{X}(n) - \mu}{\sqrt{\sigma^2/n}}$$

Consider $F_n(z)$ the cumulative distribution function of Z_n for a sample of size n .

The Central Limit Theorem states that

$$\lim_{n \rightarrow +\infty} F_n(z) = \Phi(z)$$

where $\Phi(z)$ is the cumulative distribution function of a standard Gaussian random variable (i.e., a Gaussian distribution with mean 0 and variance 1).

Given that $\lim_{n \rightarrow +\infty} S^2(n) = \sigma^2$ than, the random variable
$$\frac{\bar{X}(n) - \mu}{\sqrt{S^2(n)/n}}$$

has approximately a standard Gaussian distribution.

Analysis of the results of a simulation

For a sufficiently high value of n ,

$$P\left(-z_{1-\alpha/2} \leq \frac{\bar{X}(n) - \mu}{\sqrt{S^2(n)/n}} \leq z_{1-\alpha/2}\right) =$$
$$P\left(\bar{X}(n) - z_{1-\alpha/2}\sqrt{S^2(n)/n} \leq \mu \leq \bar{X}(n) + z_{1-\alpha/2}\sqrt{S^2(n)/n}\right) \approx 1 - \alpha$$

where $z_{1-\alpha/2}$ is the critical value of the standard Gaussian distribution ($z_{1-\alpha/2}$ is the value z such that $P(x \leq z) = 1 - \alpha/2$ where x is a random variable with a standard Gaussian distribution).

Therefore, the approximate confidence interval of $100(1-\alpha)\%$ for the average μ is given as

$$\bar{X}(n) \pm z_{1-\alpha/2}\sqrt{S^2(n)/n}$$

Analysis of the results of a simulation

The approximate confidence interval of $100(1-\alpha)\%$ for the average μ is

$$\bar{X}(n) \pm z_{1-\alpha/2} \sqrt{S^2(n)/n}$$

In MATLAB:

```
N = 20; %number of simulations
per1= zeros(1,N); %vector with N simulation values
per2= zeros(1,N); %vector with N simulation values
for it= 1:N
    [per1(it),per2(it)]= simulator();
end

alfa= 0.1; %90% confidence interval%
media = mean(per1);
term = norminv(1-alfa/2)*sqrt(var(per1)/N);
fprintf('per1 = %.2e +- %.2e\n',media,term)
media = mean(per2);
term = norminv(1-alfa/2)*sqrt(var(per2)/N);
fprintf('per2 = %.2e +- %.2e\n',media,term)
```

Analysis of the results of a simulation

The central limit theorem requires variables X_1, X_2, \dots, X_n to be independent and identically distributed (IID).

- One way to guarantee this independence is to run different simulations guaranteeing that the random values are different on the different runs.
- This is done by using different seeds on the random generators.

The confidence interval provides a measure on how close the estimated performance value might be from its real value:



- If it is small, the estimated value should be very close to the real value.
- If it is large, we don't know the quality of the estimated value.

For a given set of simulations, if the confidence interval is large, it can be made small by:

- running a larger number of simulations, and/or
- running longer simulations.

Analysis of the results of a simulation

In general, the stochastic processes have initial transient states (which are dependent on the initial conditions) before reaching the stationary state.



- In order to guarantee that the performance estimations are correct, the simulation must first warm-up to let the transient states vanish.
- If the simulated time is much higher than the warm-up time, statistical counters can be initialized at the beginning of the simulation.
- Otherwise, the statistical counters must be initialized only after the warm-up time (this time must be estimated though).